




Code Llama: Open Foundation Models for Code

...

Presented By : Mohammad mohammadi
Supervisor : Dr.Katanforosh


Introduction

- What is a large language model (LLM)?
- What is code generation?
- Why is code generation important?




Introduction

A Large Language Model (LLM) is a cutting-edge AI system adept at understanding and generating human-like text for various applications.



Introduction

Code generation is the automated creation of executable or source code from a higher-level representation, streamlining the coding process.



Introduction

Code generation is essential for automating and streamlining software development, saving time, reducing errors, and improving overall productivity by automatically creating code from higher-level representations.

Background

- Previous work on LLMs for code generation
 - The **Code2Seq** model developed by Microsoft Research was able to translate natural language descriptions of code into Python code
 - The **CodeBERT** model developed by Google AI was able to identify and complete missing code in Python functions
- Challenges of LLMs for code generation
 - Understanding Code Context
 - Generating Correct and Efficient Code
 - Biases and Fairness
 - Scalability and Efficiency
 - Explainability and Transparency

Methodology

Code Llama architecture :

- Transformer-based encoder and decoder
- Transformer-encoder with infilling
 - Filling in missing function arguments
 - Fixing syntax errors
 - Extending code functionality
- Zero-shot instruction following

Methodology

Code Llama training :

- Data Preprocessing
 - Cleans and prepares the training data
- Self-supervised Learning
 - Trains the model to predict the next token in a sequence of code
- Supervised Learning
 - Fine-tunes the model on code completion, code generation, and code summarization tasks
- Long Context and Instruction Fine-tuning
 - Fine-tunes the model on a dataset of code and natural language text that contains long contexts and natural language instructions

Methodology

Code Llama evaluation :

- HumanEval
 - A benchmark assessing a code generation model's proficiency in producing code akin to human-written code.
 - score of 53% on HumanEval
 - highest score ever achieved on HumanEval by an open-source LLM


Methodology

Code Llama evaluation :

- MBPP
 - Ability of a code generation model to write code to complete code fragments
 - Score of 55% on MBPP
 - highest score ever achieved on MBPP by an open-source LLM

The Code Llama models

- Code Llama
- Code Llama - Python
- Code Llama - Instruct



The Code Llama models

Code Llama :

- The foundational model for code generation tasks
- Trained on a dataset of over 100 billion code tokens
- Able to generate code in a variety of programming languages

The Code Llama models

Code Llama - Python :

- Specifically trained on Python code
- Accurate and efficient at generating Python code

The Code Llama models

Code Llama - Instruct :

- Fine-tuned with human instructions and self-instruct code synthesis data
- Code generation without any prior training on the specific instructions

Dataset

- Training :
 - Code Llama, from 7B to 500B tokens, focuses on publicly available code and 8% from natural language datasets.
- Dataset :
 - Predominantly near-deduplicated public code with 8% from code-related natural language datasets.
- Integration :
 - 8% of samples involve natural language discussions, enhancing model understanding
- Tokenization :
 - Byte pair encoding (BPE) aligns with Llama models, boosting MBPP performance with natural language batches.

Infilling

- Objective :
 - Predicts missing program parts for code completion, and in-code documentation generation.
- Transformation Details:
 - Training documents are character-level split into prefix, middle, and suffix parts, employing PSM and SPM formats with a 0.9 split probability.

Long context fine-tuning

- Challenges in Handling Long Sequences
- LCFT for Long Sequences:
 - Code Llama introduces Long Context Fine-Tuning (LCFT) with 16,384-token sequences, enhancing long-range capabilities without significant training cost increase.
- Extrapolation and Stability:
 - Exhibits extrapolation capabilities and stable behavior on sequences up to 100,000 tokens in experiments during fine-tuning.

Instruction fine-tuning

- Proprietary dataset
 - "RLHF V5"
 - Reinforcement learning
 - Human feedback annotation
 - Supervised Fine-Tuning
- Self-instruct
 - Focused on code-related tasks
 - Generating interview-style programming questions
- Rehearsal

Training details

Optimization :

- The optimizer used is AdamW (Loshchilov & Hutter, 2019) with β_1 and β_2 values set at 0.9 and 0.95.
- Training involves a batch size of 4M tokens presented as sequences of 4,096 tokens each

Training details

Long context fine-tuning :

- LCFT employs a learning rate of 2e-5
- Sequence length of 16,384
- Batch sizes for model sizes 7B and 13B are set to 2M tokens
- 34B model uses 1M tokens
- Training lasts for 10,000 gradient steps by default

Results

Benchmark Evaluation :

- Python benchmarks: HumanEval, MBPP, APPS
- MultiPL-E for C++, Java, PHP, C#, TypeScript, Bash
- GSM8K benchmark for mathematical reasoning

Ablation Study :

- Impact of training: scratch vs. pretrained Llama 2 model
- Evaluation of instruction fine-tuning using self-generated unit tests

Python code generation

Model Specialization:

- Massive performance gains with training on additional code-heavy tokens.
- Code Llama - Python 7B outperforms even Code Llama 13B on MBPP and HumanEval.

Scaling Specialized Models:

- Larger models consistently outperform smaller counterparts on HumanEval, MBPP.

Multilingual evaluation

Multilingual Evaluation:

- MultiPL-E benchmark
- Diverse languages: Python, C++, Java, PHP, TypeScript, C#, Bash
- Code Llama outperforms Llama 2 in all languages; Code Llama 7B surpasses Llama 2 70B

Code Llama - Python Performance:

- Comparable performance between Code Llama and Code Llama - Python

Influence of Multilingual Pre-training:

- Correlations measured between model performance on different languages
- High correlation observed among C++, C#, Java, and PHP

Infilling evaluations

Performance Cost of Infilling Training:

- Studies suggest minimal impact on test losses and slight downstream evaluation costs for multitask infilling objectives.
- Independent validation at 7B and 13B scales shows a modest decline in performance (0.6% and 1.1% loss, respectively).
- Decision to release Code Llama 7B and 13B with infilling due to its wide applicability.

Infilling evaluations

Code Infilling Benchmarks:

- Code Llama's infilling models excel in HumanEval and MultiPL-E benchmarks, outperforming models of their size.
- Superiority demonstrated in HumanEval infilling tasks and across three programming languages in the MultiPL-E benchmark.
- Notable performance increase observed when prompting models in SPM format.

Long context evaluations

Long Sequence Handling :

- Exploration of Code Llama's ability to handle long sequences through perplexity, key retrieval, and code completion tasks.
- Steady decrease in perplexity observed well beyond 16,384 tokens, showcasing stability during extrapolation.
- Strong key retrieval performance for models on the sequence length they were trained on, with slight performance decrease for longer sequences.
- LCFT (Long Context Fine-Tuning) models demonstrate significantly higher completion accuracy in single-line code completion tasks.

Long context evaluations

Performance Impact on Short Sequences:

- LCFT slightly impacts performance on short sequences in standard code synthesis benchmarks, resulting in a small decrease in scores.
- Average decrease observed on HumanEval and MBPP benchmarks.
- Despite this, the release decision is influenced by the potential benefits of handling long sequences for real-world applications.

Ablation studies

- Fine tuning Llama 2 vs. training from scratch on code
- Instruction fine-tuning
 - General helpfulness vs. coding ability
 - The value of self-instruct data
 - Unnatural model

Responsible AI and safety

Safety Benchmarks Evaluation:

- Aiming to enhance safety.
- Fine-tuning Code Llama - Instruct on Llama 2 outputs improves safety by addressing adversarial prompts.
- Safety benchmarks : evaluate truthfulness, toxicity, and bias

Responsible AI and safety

Red Teaming and Risk Evaluation:

- Red teaming exercises involve 25 Meta employees assessing risks.
- Quantitative risk evaluation on generating malicious code shows Code Llama's tendency to provide safer responses.

Responsible AI and safety

False Refusals and Safety-Performance Balance:

- Red teamers find limited evidence of false refusals, and Code Llama's safety scores skew toward the safer range.
- Safety prioritization in instruction fine-tuning may slightly degrade coding performance, maintaining a balance.

Related work

Training Data for LLMs:

- Role of Code in Training Data
- GPT-Neo and GPT-J Observations
- Open-source vs Closed-source Models
- Data Quality Impact
- Llama 2 Release and License

Related work

LLMs Landscape and Models:

- GPT-4, PaLM, Chinchilla, BLOOM, OPT
- Code-Specific LLMs: Codex, AlphaCode, GPT-NeoX, SantaCoder, StarCoder
- CodeXGlue Benchmark

Related work

Tasks Related to Code Understanding and Synthesis:

- Program Synthesis and Infilling
- Code Understanding Tasks and CodeXGlue
- Code Modifications in LLM Training
- Incorporating Program Execution in Training

Related work

Handling Long Sequences in LLMs:

- Scaling Transformers to Long Sequences
- LLMs Supporting Long Input Sequences
- Positional Information and RoPE
- Fine-tuning for Long Sequence Handling
- Challenges and Approaches in Handling Long Sequences

Discussion

Code Llama Family:

- Code Llama Model Variants (7B, 13B, 34B)
- Variants: Code Llama, Code Llama - Python, Code Llama - Instruct
- Support for Infilling and Large Contexts
- Stability in Inference Up to 100K Tokens

Discussion

Performance on Benchmarks:

- Large Context Fine-tuning and Infilling Impact
- Challenges in Standard Code Generation Benchmarks
- State-of-the-Art in Python Completion Benchmarks
- Multilingual Benchmark Performance

Discussion

Code Llama - Instruct Model:

- Zero-shot Instruction Ability
- Fine-tuning with Llama 2-Chat Distillation
- Balancing Direct Helpfulness and Safety
- Challenges in Maintaining Evaluation Scores
- Need for Improved Context and Nuance Understanding